

Learn Git and GitHub without any code!

Using the Hello World guide, you'll start a branch, write comments, and open a pull request.

[Read the guide](#)

 [IDSIA](#) / [sacred](#)

Sacred is a tool to help you configure, organize, log and reproduce experiments developed at IDSIA.

[#python](#) [#machine-learning](#) [#infrastructure](#) [#reproducible-research](#) [#reproducibility](#) [#reproducible-science](#) [#mongodb](#)

 1,279 commits

 6 branches

 0 packages

 26 releases

 72 contributors

 MIT

Branch: master ▾




[New pull request](#)














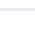


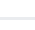
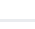
[Create new file](#)

[Upload files](#)

[Find file](#)

[Clone or download ▾](#)

 boeddeker Made git at import time optional (#724) ...	✓ Latest commit c1c19a5 on 17 Mar
 .github	Make stale time longer (#484) 12 months ago
 docs	Fix typo in docs (#715) 4 months ago
 examples	Removed create from FileStorageObserver. (#605) 10 months ago

 sacred	Made git at import time optional (#724)	3 months ago
 tests	Add option to pass a collection prefix to MongoObserver (#704)	6 months ago
 .gitignore	[API Change] Changed the way host information is collected. (#569)	10 months ago
 .pre-commit-config.yaml	Use jsonpickle handlers for numpy and pandas (#695)	7 months ago
 CONTRIBUTING.rst	More info to contribute. (#654)	9 months ago
 HISTORY.rst	Minor Patch release	6 months ago
 LICENSE.txt	added MIT license as a LICENSE.txt file	6 years ago
 MANIFEST.in	Some cleanup in the setup.py (#666)	8 months ago
 Makefile	removed unnecessary 'setup.py test' support because it was breaking a...	4 years ago
 README.rst	Link to stable version of docs in readme (#620)	9 months ago
 azure-pipelines.yml	Added a check for pre-commit in the CI and added flake8 in the pre-co...	9 months ago
 dev-requirements.txt	Update jsonpickle version in dev-requirements (#702)	6 months ago
 doc-requirements.txt	Add documentation for docstring format (#619)	9 months ago
 pyproject.toml	Use black to format code (#579)	10 months ago
 requirements.txt	Use jsonpickle handlers for numpy and pandas (#695)	7 months ago
 setup.cfg	Added a check for pre-commit in the CI and added flake8 in the pre-co...	9 months ago
 setup.py	Use jsonpickle handlers for numpy and pandas (#695)	7 months ago
 tox.ini	Update black to stable version (#688)	7 months ago

 [README.rst](#)

Sacred

*Every experiment is sacred
Every experiment is great
If an experiment is wasted
God gets quite irate*

pypi **v0.8.1** python **3.5 | 3.6 | 3.7** license **MIT** docs **passing** DOI **10.5281/zenodo.16386**

 Azure Pipelines **succeeded** coverage **85%** Scrutinizer **8.54** code style **black**

Sacred is a tool to help you configure, organize, log and reproduce experiments. It is designed to do all the tedious overhead work that you need to do around your actual experiment in order to:

- keep track of all the parameters of your experiment
- easily run your experiment for different settings
- save configurations for individual runs in a database
- reproduce your results

Sacred achieves this through the following main mechanisms:

- **ConfigScopes** A very convenient way of the local variables in a function to define the parameters your experiment uses.
- **Config Injection:** You can access all parameters of your configuration from every function. They are automatically injected by name.
- **Command-line interface:** You get a powerful command-line interface for each experiment that you can use to change parameters and run different variants.
- **Observers:** Sacred provides Observers that log all kinds of information about your experiment, its dependencies, the configuration you used, the machine it is run on, and of course the result. These can be saved to a MongoDB, for easy

access later.

- **Automatic seeding** helps controlling the randomness in your experiments, such that the results remain reproducible.

Example

Script to train an SVM on the iris dataset	The same script as a Sacred experiment
<pre>from numpy.random import permutation from sklearn import svm, datasets C = 1.0 gamma = 0.7 iris = datasets.load_iris() perm = permutation(iris.target.size) iris.data = iris.data[perm] iris.target = iris.target[perm] clf = svm.SVC(C, 'rbf', gamma=gamma) clf.fit(iris.data[:90], iris.target[:90]) print(clf.score(iris.data[90:], iris.target[90:])))</pre>	<pre>from numpy.random import permutation from sklearn import svm, datasets from sacred import Experiment ex = Experiment('iris_rbf_svm') @ex.config def cfg(): C = 1.0 gamma = 0.7 @ex.automain def run(C, gamma): iris = datasets.load_iris() per = permutation(iris.target.size) iris.data = iris.data[per] iris.target = iris.target[per] clf = svm.SVC(C, 'rbf', gamma=gamma) clf.fit(iris.data[:90], iris.target[:90]) return clf.score(iris.data[90:], iris.target[90:])))</pre>

Documentation

The documentation is hosted at [ReadTheDocs](#).

Installing

You can directly install it from the Python Package Index with pip:

```
pip install sacred
```

Or if you want to do it manually you can checkout the current version from git and install it yourself:

```
git clone https://github.com/IDSIA/sacred.git
cd sacred
python setup.py install
```

You might want to also install the `numpy` and the `pymongo` packages. They are optional dependencies but they offer some cool features:

```
pip install numpy, pymongo
```

Tests

The tests for sacred use the [pytest](#) package. You can execute them by running `pytest` in the sacred directory like this:

```
pytest
```

There is also a config file for [tox](#) so you can automatically run the tests for various python versions like this:

```
tox
```

Contributing

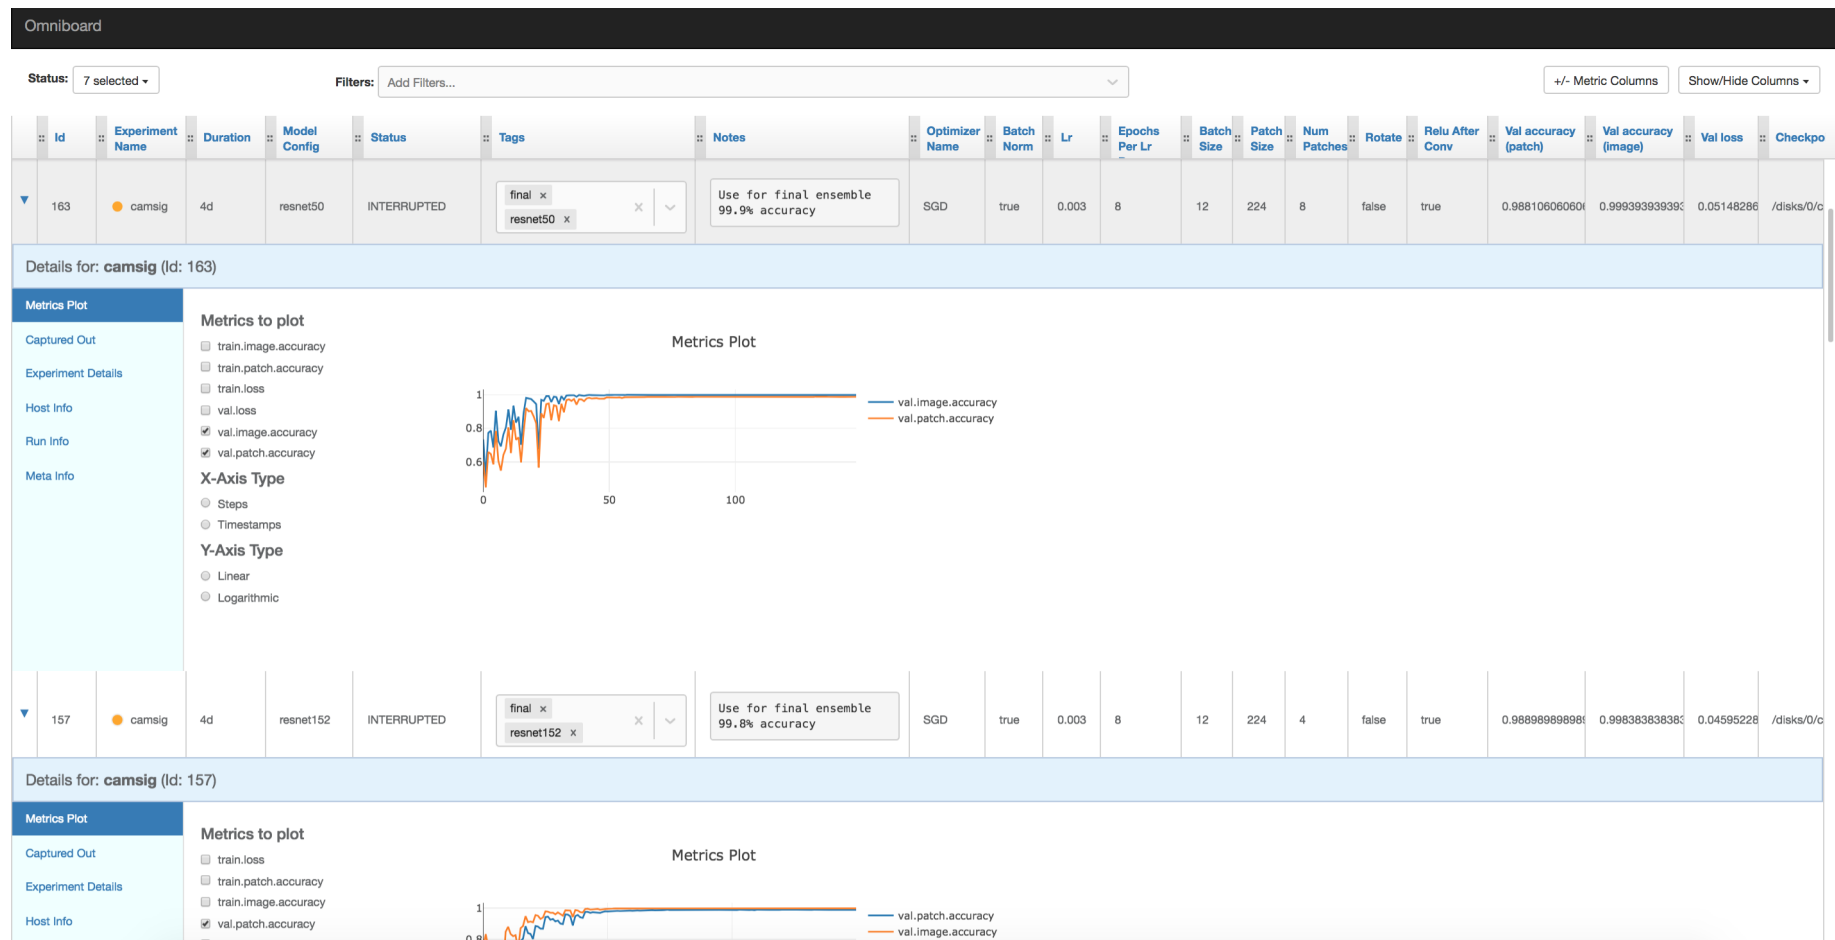
If you find a bug, have a feature request or want to discuss something general you are welcome to open an [issue](#). If you have a specific question related to the usage of sacred, please ask a question on StackOverflow under the [python-sacred tag](#). We value documentation a lot. If you find something that should be included in the documentation please document it or let us know whats missing. If you are using Sacred in one of your projects and want to share your code with others, put your repo in the Projects using Sacred <docs/projects_using_sacred.rst>_list. Pull requests are highly welcome!

Frontends

At this point there are three frontends to the database entries created by sacred (that I'm aware of). They are developed externally as separate projects.

Omniboard


Omniboard																			
Status: 7 selected ▾		Filters: Add Filters... ▾										+/- Metric Columns			Show/Hide Columns ▾				
Id	Experiment Name	Duration	Model Config	Status	Tags	Notes	Optimizer Name	Batch Norm	Lr	Epochs Per Lr	Batch Size	Patch Size	Num Patches	Rotate	Relu After Conv	Val accuracy (patch)	Val accuracy (image)	Val loss	Checkpoint
166	● camsig	2d	densenet161	PROBABLY_DEAD	ignore x ▾	Ran out of time	SGD	true	0.003	8	8	224	4	false	true	0.86671810699	0.91982167352	0.42694607	/disks/0/c
163	● camsig	4d	resnet50	INTERRUPTED	final x resnet50 x ▾	Use for final ensemble 99.9% accuracy	SGD	true	0.003	8	12	224	8	false	true	0.98810606060	0.99939393939	0.05148286	/disks/0/c
157	● camsig	4d	resnet152	INTERRUPTED	final x resnet152 x ▾	Use for final ensemble 99.8% accuracy	SGD	true	0.003	8	12	224	4	false	true	0.98898989898	0.99838383838	0.04595228	/disks/0/c
130	● camsig	2d	vgg11	INTERRUPTED	ignore x ▾	66,954 trainable params Bad accuracy/loss after 67 epochs Giving up	SGD	true	0.003	8	12	224	10		true	0.51474747474	0.74565656565	1.44083583	/disks/0/c
129	● camsig	3d	resnet152	INTERRUPTED	ignore x ▾	263,562 trainable params	SGD	true	0.003	8	12	224	4		true	0.73777777777	0.86262626262	0.84879525	/disks/0/c
127	● camsig	3d	vgg16_bn	INTERRUPTED	ignore x ▾	Enter Notes	SGD	true	0.003	8	12	224	6		true	0.66114478114	0.84424242424	1.06450540	/disks/0/c
123	● camsig	16h	B	COMPLETED	candidate x ▾	Enter Notes	SGD	true	0.015	10	12	64	128		true	0.90258522727	0.98727272727	0.36969404	/disks/0/c
122	● camsig	16h	B	COMPLETED	candidate x ▾	Enter Notes	SGD	true	0.015	10	12	64	128		true	0.87825284090	0.98727272727	0.44598068	/disks/0/c
121	● camsig	12h	B	COMPLETED	candidate x ▾	Enter Notes	SGD	true	0.015	10	12	64	128		true	0.89066761363	0.99272727272	0.40099986	/disks/0/c



Omniboard is a web dashboard that helps in visualizing the experiments and metrics / logs collected by sacred. Omniboard is written with React, Node.js, Express and Bootstrap.

Incense



```
[16]: exp.artifacts['confusion_matrix'].show(figsize=(10, 10));
```



```
[17]: exp.artifacts['confusion_matrix'].save()
```

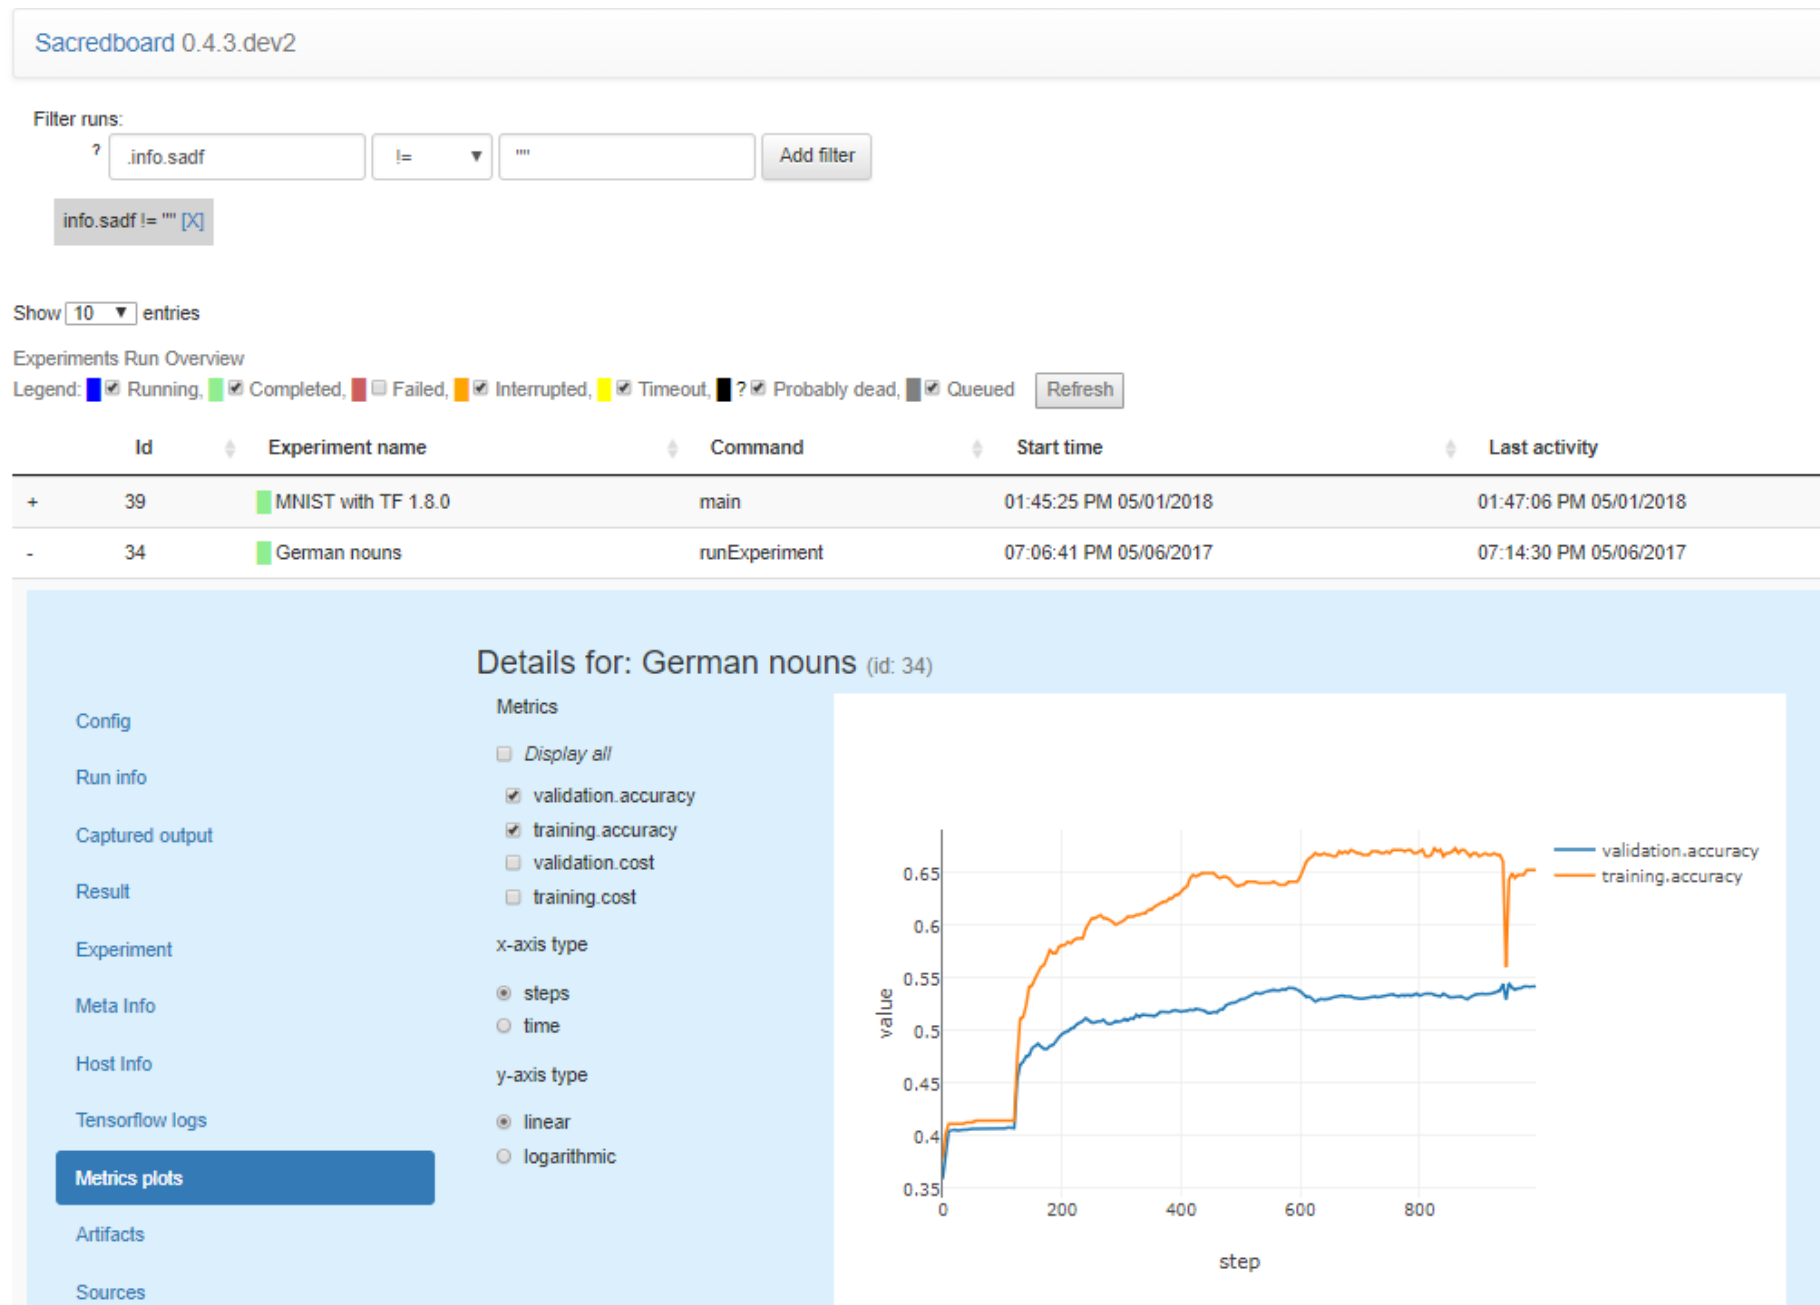
```
[23]: exp.metrics['training_acc'].plot()  
exp.metrics['training_loss'].plot()  
plt.legend()
```

```
[23]: <matplotlib.legend.Legend at 0x7f5b34d9feb8>
```



Incense is a Python library to retrieve runs stored in a MongoDB and interactively display metrics and artifacts in Jupyter notebooks.

Sacredboard



Sacredboard is a web-based dashboard interface to the sacred runs stored in a MongoDB.

Neptune



Short ID	Owner	Tags	valid_auc	train_auc	features_...	train_spli...	valid_spli...
CRED-93	kamil	lgbm features_v1	0.773536	0.807047	e1f5a473fe...	41bf231b7...	4b61d05c0...
CRED-92	kamil	lgbm features_v1	0.765647	0.790213	e1f5a473fe...	41bf231b7...	4b61d05c0...
CRED-91	kamil	lgbm features_v1	0.767254	0.78022	e1f5a473fe...	41bf231b7...	4b61d05c0...
CRED-89	kamil	lgbm features_v1	0.772704	0.788251	e1f5a473fe...	41bf231b7...	4b61d05c0...

The screenshot displays the Neptune ML web interface. At the top, a header bar shows the project name 'CRED-85', the user 'jakub-czakon', and several tags: 'lgbm', 'features_v1', and two numerical values '0.775187' and '0.797085' highlighted in green. Below this is a navigation bar with tabs for 'Projects', 'Explore', 'Wiki', 'Notebooks' (6), 'Experiments' (30), and 'Settings'. A user profile dropdown for 'jakub-czakon' is on the right.

The main content area shows a project page for '2.0-modeling_lgbm' with the status 'All changes saved'. The page title is 'Best Model' followed by a green 'S' icon and 'CRED-83'. Below the title, there are links for 'metric' (CRED-83/channels) and 'hyperparameters' (CRED-83/details). The metric value is '0.775489'.

A sidebar on the left lists pages: '0.0-report', '0.1-open_problems', '0.2-onboarding', '1.0-feature_extraction_v0', '1.1-feature_extraction_v1', and '2.0-modeling_lgbm' (selected).

A 'Parameters' table is shown, listing three parameters:

#	Parameter	Value	Description
1	num_boost_round	1000	
2	early_stopping_rounds	100	
3	metric	auc	

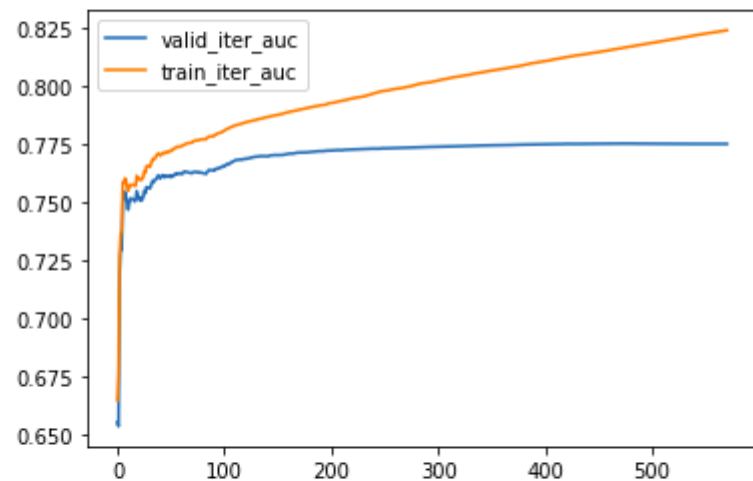
On the right, a comment box shows a comment from 'jakub-czakon' dated '2019/08/13 11:27' with the text 'What do you think kamil?'. There are 'Save' and 'Cancel' buttons, and a 'Resolve' button.

Neptune is a web service that lets you visualize, organize and compare your experiment runs. Once things are logged to Neptune you can share it with others, add comments and even access objects via experiment API:

```
[3]: exp = project.get_experiments(id=['CRED-80'])[0]
     exp
```

```
[3]: Experiment(CRED-80)
```

```
[4]: fig, ax = plt.subplots()
     df = exp.get_numeric_channels_values('train_iter_auc', 'valid_iter_auc')
     ax.plot(df.x, df.valid_iter_auc, label='valid_iter_auc')
     ax.plot(df.x, df.train_iter_auc, label='train_iter_auc')
     plt.legend()
     plt.show()
```



```
[5]: exp.get_properties()
```

```
[5]: {'features_path': 'data/processed/features_joined_v1.csv',
      'features_version': 'elf5a473fedb57d12a792c02cc52993d',
      'train_split_version': '41bf231b7dc76abb7a2fe15e5cbd9d20',
      'valid_split_version': '4b61d05c05aa16f71015614e3e3b1d1a'}
```

In order to log your runs to Neptune, all you need to do is add an observer:

```
from neptunecontrib.monitoring.sacred import NeptuneObserver
ex.observers.append(NeptuneObserver(api_token='YOUR_API_TOKEN',
```

```
project_name='USER_NAME/PROJECT_NAME' ) )
```

For more info, check the [neptune-contrib](#) library.

SacredBrowser

Database

- INPUT
 - study1
 - admin
 - assistant
 - binding
 - random_search
 - binding_via_rc
 - caidb

em
training
verbose
network_spec
dataset
seed
net_filename

UP
DOWN

Instructions: Enter several lines with conditions. The basic form is
ConfigParam : value. The value is automatically converted to int, float, or string.
An item is displayed if the conditions on all lines are fulfilled,
alternative ("or") conditions can be written in list style:
ConfigParam: [val1, val2, etc]

☐ Allow delete without confirmation

Result view mode ☒ Raw ☐ Rounded ☐ Percent

Sort Dialog

Delete Copy Full entry New search Clear Undo

	em	training	verbose	network_spec	dataset	seed	net_filename	Result 1
0			false	Fr100		939888202	Networks/...	0.77072582...
1			false	Ft500		327278722	Networks/...	0.28177538...
2			false	Ft1000		312145402	Networks/...	0.02424216...
3			false	Ft100		875640577	Networks/...	0.42828769...
4			false	Fr100		749643037	Networks/...	0.46656010...
5			false	Ft100		691571401	Networks/...	0.00299657...
6			false	Fr500		848455680	Networks/...	0.02332285...
7			false	Fr1000		509098946	Networks/...	---

Connect to MongoDB instance Averages (excluding duplicates): 0.32 Reset column widths

Loaded 604 entries, found 0 possible duplicates, 3 without result

SacredBrowser is a PyQt4 application to browse the MongoDB entries created by sacred experiments. Features include custom queries, sorting of the results, access to the stored source-code, and many more. No installation is required and it can connect to a local database or over the network.

Prophet

Prophet is an early prototype of a webinterface to the MongoDB entries created by sacred experiments, that is discontinued. It requires you to run [RestHeart](#) to access the database.

Related Projects

Sumatra

Sumatra is a tool for managing and tracking projects based on numerical simulation and/or analysis, with the aim of supporting reproducible research. It can be thought of as an automated electronic lab notebook for computational projects.

Sumatra takes a different approach by providing commandline tools to initialize a project and then run arbitrary code (not just python). It tracks information about all runs in a SQL database and even provides a nice browser tool. It integrates less tightly with the code to be run, which makes it easily applicable to non-python experiments. But that also means it requires more setup for each experiment and configuration needs to be done using files. Use this project if you need to run non-python experiments, or are ok with the additional setup/configuration overhead.

Future Gadget Laboratory

FGLab is a machine learning dashboard, designed to make prototyping experiments easier. Experiment details and results are sent to a database, which allows analytics to be performed after their completion. The server is FGLab, and the clients are FGMachines.

Similar to Sumatra, FGLab is an external tool that can keep track of runs from any program. Projects are configured via a JSON schema and the program needs to accept these configurations via command-line options. FGLab also takes the role of a basic scheduler by distributing runs over several machines.

CDE

By tracing system calls during program execution CDE creates a snapshot of **all** used files and libraries to guarantee the ability to reproduce any unix program execution. It *only* solves reproducibility, but it does so thoroughly.

License

This project is released under the terms of the [MIT license](#).

Citing Sacred

K. Greff, A. Klein, M. Chovanec, F. Hutter, and J. Schmidhuber, 'The Sacred Infrastructure for Computational Research', in Proceedings of the 15th Python in Science Conference (SciPy 2017), Austin, Texas, 2017, pp. 49–56.